

# Concurrent Programming Notes

v0.015 – November 21, 2022

Eduardo Bonelli



# Preface

These course notes provide supporting material for CS511.



# Contents

<b>Preface</b>	<b>i</b>
<b>1 Shared Memory Model and Transition Systems</b>	<b>1</b>
1.1 Shared Memory Model . . . . .	1
1.2 Transition Systems . . . . .	3
1.3 Atomicity . . . . .	6
1.4 The Mutual Exclusion Problem . . . . .	8
<b>2 Semaphores</b>	<b>9</b>
2.1 Introduction . . . . .	9
2.2 The MEP Problem Revisited . . . . .	9
2.3 More Examples . . . . .	11
2.3.1 Thread Dumps . . . . .	12
2.4 Classical Synchronization Problems . . . . .	14
2.4.1 Producers/Consumers . . . . .	14
2.4.2 Readers/Writers . . . . .	16
2.4.3 Barrier Synchronization . . . . .	16
<b>3 Monitors</b>	<b>21</b>
3.1 A monitor implementing a semaphore . . . . .	22
3.2 Producers/Consumers . . . . .	23
3.3 Readers/Writers . . . . .	24
<b>4 Promela</b>	<b>29</b>
4.1 Syntax . . . . .	29
4.1.1 Examples involving Loops . . . . .	30
4.1.2 Expressions as blocking commands . . . . .	30
4.1.3 Macros . . . . .	32
4.2 Assertion-Based Model Checking . . . . .	35
4.2.1 The Bar Problem Revisited . . . . .	35
4.2.2 The MEP Problem . . . . .	38
4.3 Non-Progress Cycles . . . . .	40
4.3.1 The Zoo Problem Revisited . . . . .	44
<b>5 Solution to Selected Exercises</b>	<b>47</b>



# Chapter 1

## Shared Memory Model and Transition Systems

This chapter...

### 1.1 Shared Memory Model

We begin with an example of a program in Groovy.

```
1  int x = 0
3  Thread.start { //P
    x = 1
5  }
7  Thread.start { //Q
    x = 2
9  }
```

ex1.groovy

This program declares a shared variable `x`, sets it to 0 and then spawns two threads. The first thread sets `x` to 1 and the second to 2. After this program terminates, the value of `x` may either be 1 or 2. The variable `x` is said to be shared in the sense that it is visible to (or its scope includes) both threads<sup>1</sup>.



Semicolons are optional in Groovy

Assuming this program is stored in a file called `ex1.groovy`, it may be executed using the terminal as follows:

---

<sup>1</sup>From the point of view of Groovy it is actually a local variable. In Groovy, global variables are declared by omitting the type annotation.

```
$ groovy ex1
2 $
```

bash

Since our program contains no output statements, there is no visible effect from its execution. The following example, waits for P and Q to terminate using the built-in method `join` and then prints the value of `x`:

```
int x = 0
2
3 P = Thread.start { //P
4     x = 1
5 }
6
7 Q = Thread.start { //Q
8     x = 2
9 }
10
11 P.join() // Wait for P to terminate
12 Q.join() // Wait for Q to terminate
println x
```

ex2.groovy

Assuming this program is stored in a file called `ex2.groovy`, it may be executed using the terminal as follows:

```
$ groovy ex2
2 2
3 $
```

bash

Repeated execution will most likely produce 2 since P is spawned before Q and runs immediately. It is entirely possible, however, to obtain 1 as a result.

The following example is a Groovy program that prints characters.

```
1 Thread.start { //P
2     print "A"
3     print "B"
4 }
5
6 Thread.start { //Q
7     print "C"
8 }
```

What are the possible outputs one may obtain from executing it? It can print three possible sequences of characters, namely ABC, ACB, CAB. What about the following program?

```
2 Thread.start { //P
3     print "A"
4     print "B"
5 }
6
7 Thread.start { //Q
8     print "C"
9     print "D"
```



```
}

```

Clearly the number of possible executions, also called interleavings, grows exponentially with the number of instructions in each thread. Indeed, if P has  $m$  instructions and Q has  $n$  instructions, then there are

$$\binom{m+n}{m} = \frac{(m+n)!}{m!n!}$$

This makes it difficult to reason about concurrent programs: there are simply too many interleavings to consider; we never know whether one such interleaving might lead our code to produce an unwanted result. We clearly need some rigorous device to be able to model all such possible interleavings and check whether they satisfy our intended properties. A device that describes the run-time execution of a concurrent program. There is a further, equally important reason, why we need this device. Consider the following program:

```
1 int x=0 // shared variable
3 P = Thread.start {
    x = x+1
5 }
Q = Thread.start {
    x = x+1
7 }
9
P.join() // wait for P to terminate
11 Q.join() // wait for Q to terminate
13 println(x)
```

Its execution produces 1 as output!

```
1 $ groovy ex1
1
3 $
```

bash

How is that possible?

## 1.2 Transition Systems

This section introduces transition systems, a device we use to model the run-time behavior of concurrent programs. After defining transition systems, we illustrate how to associate a transition system to Groovy programs. By doing so, we assign “meaning” to our concurrent programs. It should be mentioned that we will associate transition systems only to a subset of simple Groovy programs, not arbitrary ones.

A **Transition System**  $\mathcal{A}$  is a tuple  $(S, \rightarrow, I)$  where

- $S$  is a set of states;
- $\rightarrow \subseteq S \times S$  is a transition relation; and
- $I \subseteq S$  is a set of initial states.

We say that  $\mathcal{A}$  is finite if  $S$  is finite. Also, we write  $s \rightarrow s'$  for  $(s, s') \in \rightarrow$ .


We illustrate, in this first example, how to model the runtime execution of Example 1.1, repeated below:

```

1  int x = 0
3  Thread.start { //P
    x = 1
5  }
7  Thread.start { //Q
    x = 2
9  }

```

The states of our transition system will consist of 3-tuples containing the instruction pointer for  $p$ , the instruction pointer for  $q$  and the value of  $x$ . The initial state is signalled with a small arrow. The hyphen indicates that there are no further instructions to be executed by that thread.

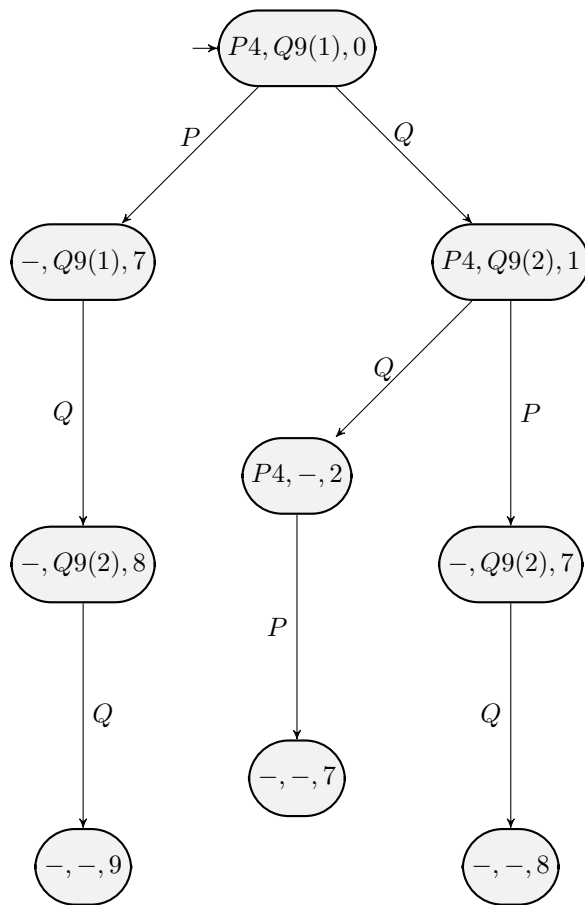


How we hardcode “for”-loops with a constant upper bound.

```

1  int x = 0
2
3  Thread.start { //P
4      x = 7
5  }
6
7  Thread.start { //Q
8      2.times {
9          x = x+1
10     }
11 }

```



How we distinguish local variables with the same name using "local\_P" and "local\_Q" in the state format.

```

1  int x = 0 // shared variable
3  Thread.start { //P
    int local = x
5   x = local+1 // atomic
6  }
7
8  Thread.start { //Q
9   int local = x
10  x = local+1 // atomic
11 }
  
```

How we deal with "while"-loops.

```

1  int x=0 // shared variable
3  Thread.start { //P
    while (x<1) {
5     print x
    }
  }
  
```

```

7 }
9 Thread.start { //Q
    x = x + 1
11 }

```

How we decorate transitions with the output string of a print

```

1 int x = 0 // shared variable
3 Thread.start { //P
    x = x + 1
5     x = x + 1
    }
7
9 Thread.start { //Q
    while (x >= 2)
        print x
11 }

```

```

1 int counter=0 // shared variable
3 P = Thread.start {
    50.times {
5         counter = counter+1
        }
7     }
9 Q = Thread.start {
    50.times {
        counter= counter+1
11     }
    }
13
15 P.join() // wait for P to finish
    Q.join() // wait for Q to finish
17
19 println counter // print value of counter

```

## 1.3 Atomicity

Consider the following program:

```

1 x=0
3 Thread.start { //P
    x = x + 1
5     println x
    }
7
9 Thread.start { //Q
    x = x + 1
11     println x
    }

```

One would expect 1 and 2, or 2 and 2 to be printed. These are indeed possible outputs. However, 1 and 1 is also possible:

```

1 $ groovy ex3
1
3 1
$

```

bash

The reason is that assignment is not an atomic operation, rather it is decomposed into more fine grained (bytecode) operations. It is the latter that are interleaved. Let's take a closer look at those fine grained operations. Consider the following Java class that spawn two threads, each of which updates a shared variable:

```

class A implements Runnable {
2
    static int x=0;
4
    public void run() {
6
        x=x+1;
    }
8
    public static void main(String[] args) {
10
        new Thread(new A()).start();
        new Thread(new A()).start();
12
    }
}

```

A.java

We compile it and look at the resulting bytecode by using `javap`, the Java class file disassembler:

```

$ javac A.java
$ javap -c A
Compiled from "A.java"
4 class A implements java.lang.Runnable {
    static int x;
6
    A();
8    Code:
        0: aload_0
        1: invokespecial #1                  // Method java/lang/Object."<init>":()V
        4: return
12
    public void run();
14    Code:
        0: getstatic     #7                  // Field x:I
        3: iconst_1
        4: iadd
        5: putstatic     #7                  // Field x:I
        8: return
16
    public static void main(java.lang.String[]);
20    Code:
        0: new           #13                 // class java/lang/Thread
        3: dup
22

```

```

26      4: new          #8          // class A
      7: dup
      8: invokespecial #15          // Method "<init>":()V
28     11: invokespecial #16          // Method java/lang/Thread."<init>":(Ljava
      14: invokevirtual #19          // Method java/lang/Thread.start:()V
30     17: new          #13          // class java/lang/Thread
      20: dup
32     21: new          #8          // class A
      24: dup
34     25: invokespecial #15          // Method "<init>":()V
      28: invokespecial #16          // Method java/lang/Thread."<init>":(Ljava
36     31: invokevirtual #19          // Method java/lang/Thread.start:()V
      34: return
38
40     static {};
      Code:
      0: iconst_0
42     1: putstatic      #7          // Field x:I
      4: return
44 }

```

bash

The only lines we are interested are lines 15 to 18. Each thread has a JVM stack. Every time a method is called, a new frame is created (heap-allocated) and stored on the JVM stack for that thread. Each frame has its own array of local variables, its own operand stack, and a reference to the run-time constant pool of the class of the current method. The instruction  $x=x+1$  is compiled to four bytecode instructions whose meaning can be read off from their opcodes:

```

2      0: getstatic      #7          // Field x:I
      3: iconst_1
      4: iadd
4      5: putstatic      #7          // Field x:I

```

It is these operations, for each thread, that get interleaved. Thus, it is possible to have the following interleaving:

```

2      0(P): getstatic      #7          // Field x:I
      0(Q): getstatic      #7          // Field x:I
      3(P): iconst_1
4      3(Q): iconst_1
      4(P): iadd
6      4(Q): iadd
      5(P): putstatic      #7          // Field x:I
8      5(Q): putstatic      #7          // Field x:I

```

These instructions end up storing 1 in  $x$ .

## 1.4 The Mutual Exclusion Problem

# Chapter 2

## Semaphores

### 2.1 Introduction

### 2.2 The MEP Problem Revisited

Consider the following solution to the MEP problem using a binary semaphore presented in listing 2.1<sup>1</sup>.


```
1 Semaphore mutex= new Semaphore(1)
2
3 Thread.start { //P
4     while (true) {
5         mutex.acquire()
6         mutex.release()
7     }
8 }
9 Thread.start { //Q
10    while (true) {
11        mutex.acquire()
12        mutex.release()
13    }
14 }
```

**Listing 2.1:** Solution to MEP using a binary semaphore

One easy way to verify that all three properties of MEP are upheld is to construct its transition system and then analyze these properties. This requires a means for representing semaphores. Since a semaphore is an object with state and the latter includes the number of permits and the set of blocked processes, we shall model `mutex` using the expression `mutex[i,s]` where `i` is the number of permits and `s` is a set of blocked processes. Moreover, we use the “!” symbol as instruction pointer in the states of our transition systems to indicate that there are no instructions ready to execute. For example, a state such as  $P6,!,mutex[0,\{Q11\}]$ , reflects that only P can be scheduled for execution, there are no permits available in `mutex` and one thread is blocked on

---

<sup>1</sup>Groovy requires that you import the Semaphore class in order to be able to use it. All code excerpts involving semaphores should thus include, at the top, the line `import java.util.concurrent.Semaphore`. This is typically omitted in our examples.



**Figure 2.1:** Transition System for the solution to MEP using a binary semaphore

`mutex` waiting for a permit to become available, namely `Q`. Figure ?? is the transition system for the listing in Figure 2.1.

Consider the setting where the above solution is applied to three threads wanting to access their CS. This is illustrated in Listing 2.2. Although `Mutex` and `Absence of Livelock` are upheld, `Freedom From Starvation` is not. Indeed, consider the scenario where `P` goes in and `Q` and `R` try to get in and are both blocked and placed in the set of blocked processes for `mutex`. [COMPLETE]??

This is easily solved by having the set of blocked processes in `mutex` be a queue. Such semaphores are called fair semaphores. This is achieved by using an alternative constructor for semaphores that includes a fairness parameter

```
Semaphore(int permits, boolean fair)
```

Replacing line 1 in Listing 2.2 with `Semaphore mutex= new Semaphore(1,true)` suffices to obtain a correct solution to the MEP for any number of threads.

```

1 Semaphore mutex= new Semaphore(1)
2
3 Thread.start { //P
4     while (true) {
5         mutex.acquire()
6         mutex.release()
7     }
8 }
9 Thread.start { //Q
10    while (true) {
11        mutex.acquire()
12        mutex.release()
13    }

```



```

14 }
15 Thread.start { //R
16     while (true) {
17         mutex.acquire()
18         mutex.release()
19     }
20 }

```

**Listing 2.2:** Attempt at solving the MEP using a binary semaphore for  $N=3$

## 2.3 More Examples

```

Thread.start {
2     println "A"
    println "B"
4 }
Thread.start {
6     println "C"
    println "D"
8 }

```

```

Semaphore cAfterA = new Semaphore(0)
2
Thread.start {
4     println "A"
    mutex.release()
6     println "B"
}
Thread.start {
8     mutex.acquire()
10    println "C"
    println "D"
12 }

```

Consider the following example which prints any (infinite) sequence of “a”s and “b”s:

```

Thread.start { //P
2     while (true) {
        print "a"
4     }
}
6
Thread.start { //Q
8     while (true) {
        print "b"
10    }
}

```

Using semaphores, how would you ensure that only the infinite sequence “aabaabaab...” is printed? Hint: make use of two semaphores, `a` and `b`, enabling the execution of an iteration in `P` and an iteration in `Q`, respectively.

Here is a solution.

```

import java.util.concurrent.Semaphore
2

```

```

Semaphore a = new Semaphore(2)
4 Semaphore b = new Semaphore(0)

6 Thread.start { //P
    while (true) {
8         a.acquire()
        print "a"
10        b.release()
    }
12 }

14 Thread.start { //Q
    while (true) {
16        b.acquire(2)
        print "b"
18        a.release(2)
    }
20 }

```

### 2.3.1 Thread Dumps

We can check the current thread dump of the our Groovy/Java application as follows. Let's use the example above. First we modify our code so that we give our threads an easy to spot name and remove the lines that print. The result is below; we'll call it `ex1.groovy`.

```

import java.util.concurrent.Semaphore

2 Semaphore a = new Semaphore(2)
3 Semaphore b = new Semaphore(0)

6 Thread.start { //P
    Thread.currentThread().setName("P Thread");
8    while (true) {
        a.acquire()
10        // print "a"
        b.release()
    }
12 }

14 Thread.start { //Q
16    Thread.currentThread().setName("Q Thread");
    while (true) {
18        b.acquire()
        b.acquire()
20        // print "b"
        a.release()
22        a.release()
    }
24 }

```

Now we run it in the background, use `jstack` to obtain the stack trace of the bash job and send the output to a text file `thead-dump.txt`

```

$ groovy ex1 &
2 [1] 23275
$ jstack -l 23275 > thread-dump.txt

```

```
4 $ kill %1
    [1] + 23275 exit 143    groovy ex1
6 $ emacs thread-dump.txt
```

```
bash
```

The dump contains information on all threads involved in our application. We'll just show an excerpt that mentions `p` and `q`. We can see that the former is in a `RUNNABLE` state and the latter is in a `WAITING` state. We can also see the current instruction being executed in each thread.

```

1  "P Thread" #17 prio=5 os_prio=31 cpu=5910.73ms elapsed=10.91s tid=0x00007f7d9412b400 nid=27655 runnable
2  java.lang.Thread.State: RUNNABLE
3      at jdk.internal.misc.Unsafe.unpark(java.base@18.0.1.1/Native Method)
4      at java.util.concurrent.locks.LockSupport.unpark(java.base@18.0.1.1/LockSupport.java:177)
5      at java.util.concurrent.locks.AbstractQueuedSynchronizer.signalNext(java.base@18.0.1.1/AbstractQueuedSynchronizer.java:432)
6      at java.util.concurrent.locks.AbstractQueuedSynchronizer.releaseShared(java.base@18.0.1.1/AbstractQueuedSynchronizer.java:432)
7      at java.util.concurrent.Semaphore.release(java.base@18.0.1.1/Semaphore.java:432)
8      at java.lang.invoke.LambdaForm$DMH/0x0000000800d28000.invokeVirtual(java.base@18.0.1.1/LambdaForm$DMH/0x0000000800e32c00.invoke(java.base@18.0.1.1/LambdaForm$MH)
9      at java.lang.invoke.LambdaForm$MH/0x0000000800e2b400.guardWithCatch(java.base@18.0.1.1/LambdaForm$MH/0x0000000800e27800.guard(java.base@18.0.1.1/LambdaForm$MH)
10     at java.lang.invoke.DelegatingMethodHandle$Holder.delegate(java.base@18.0.1.1/DelegatingMethodHandle$Holder.java:577)
11     at java.lang.invoke.LambdaForm$MH/0x0000000800e27800.guard(java.base@18.0.1.1/LambdaForm$MH)
12     at java.lang.invoke.DelegatingMethodHandle$Holder.delegate(java.base@18.0.1.1/DelegatingMethodHandle$Holder.java:577)
13     at java.lang.invoke.LambdaForm$MH/0x0000000800e27800.guard(java.base@18.0.1.1/LambdaForm$MH)
14     at java.lang.invoke.Invokers$Holder.linkToCallSite(java.base@18.0.1.1/Invokers$Holder)
15     at ex1$_run_closure1.doCall(ex1.groovy:12)
16     at ex1$_run_closure1.doCall(ex1.groovy)
17     at java.lang.invoke.DirectMethodHandle$Holder.invokeSpecial(java.base@18.0.1.1/DirectMethodHandle$Holder.java:577)
18     at java.lang.invoke.LambdaForm$MH/0x0000000800c1c800.invoke(java.base@18.0.1.1/LambdaForm$MH)
19     at java.lang.invoke.Invokers$Holder.invokeExact_MT(java.base@18.0.1.1/Invokers$Holder)
20     at jdk.internal.reflect.DirectMethodHandleAccessor.invokeImpl(java.base@18.0.1.1/DirectMethodHandleAccessor.java:577)
21     at jdk.internal.reflect.DirectMethodHandleAccessor.invoke(java.base@18.0.1.1/DirectMethodHandleAccessor.java:577)
22     at java.lang.reflect.Method.invoke(java.base@18.0.1.1/Method.java:577)
23     at org.codehaus.groovy.reflection.CachedMethod.invoke(CachedMethod.java:343)
24     at groovy.lang.MetaMethod.doMethodInvoke(MetaMethod.java:328)
25     at org.codehaus.groovy.runtime.metaclass.ClosureMetaClass.invokeMethod(ClosureMetaClass.java:279)
26     at groovy.lang.MetaClassImpl.invokeMethod(MetaClassImpl.java:1009)
27     at groovy.lang.Closure.call(Closure.java:418)
28     at groovy.lang.Closure.call(Closure.java:412)
29     at groovy.lang.Closure.run(Closure.java:500)
30     at java.lang.Thread.run(java.base@18.0.1.1/Thread.java:833)
31
32  Locked ownable synchronizers:
33      - None
34
35  "Q Thread" #18 prio=5 os_prio=31 cpu=6110.53ms elapsed=10.91s tid=0x00007f7d9411fa00 nid=28163 runnable
36  java.lang.Thread.State: WAITING (parking)
37      at jdk.internal.misc.Unsafe.park(java.base@18.0.1.1/Native Method)
38      - parking to wait for <0x00000006180b9960> (a java.util.concurrent.Semaphore$NonfairSync)
39      at java.util.concurrent.locks.LockSupport.park(java.base@18.0.1.1/LockSupport.java:211)
40      at java.util.concurrent.locks.AbstractQueuedSynchronizer.acquire(java.base@18.0.1.1/AbstractQueuedSynchronizer.java:432)
41      at java.util.concurrent.locks.AbstractQueuedSynchronizer.acquireSharedInterruptibly(java.base@18.0.1.1/AbstractQueuedSynchronizer.java:432)
42      at java.util.concurrent.Semaphore.acquire(java.base@18.0.1.1/Semaphore.java:318)
43      at java.lang.invoke.LambdaForm$DMH/0x0000000800d28000.invokeVirtual(java.base@18.0.1.1/LambdaForm$DMH/0x0000000800e32c00.invoke(java.base@18.0.1.1/LambdaForm$MH)
44      at java.lang.invoke.LambdaForm$MH/0x0000000800e2b400.guardWithCatch(java.base@18.0.1.1/LambdaForm$MH/0x0000000800e27800.guard(java.base@18.0.1.1/LambdaForm$MH)
45      at java.lang.invoke.DelegatingMethodHandle$Holder.delegate(java.base@18.0.1.1/DelegatingMethodHandle$Holder.java:577)
46      at java.lang.invoke.LambdaForm$MH/0x0000000800e27800.guard(java.base@18.0.1.1/LambdaForm$MH)
47      at java.lang.invoke.Invokers$Holder.linkToCallSite(java.base@18.0.1.1/Invokers$Holder)
48      at ex2$_run_closure1.doCall(ex2.groovy:12)
49      at ex2$_run_closure1.doCall(ex2.groovy)
50      at java.lang.invoke.DirectMethodHandle$Holder.invokeSpecial(java.base@18.0.1.1/DirectMethodHandle$Holder.java:577)
51      at java.lang.invoke.LambdaForm$MH/0x0000000800c1c800.invoke(java.base@18.0.1.1/LambdaForm$MH)
52      at java.lang.invoke.Invokers$Holder.invokeExact_MT(java.base@18.0.1.1/Invokers$Holder)
53      at jdk.internal.reflect.DirectMethodHandleAccessor.invokeImpl(java.base@18.0.1.1/DirectMethodHandleAccessor.java:577)
54      at jdk.internal.reflect.DirectMethodHandleAccessor.invoke(java.base@18.0.1.1/DirectMethodHandleAccessor.java:577)
55      at java.lang.reflect.Method.invoke(java.base@18.0.1.1/Method.java:577)
56      at org.codehaus.groovy.reflection.CachedMethod.invoke(CachedMethod.java:343)
57      at groovy.lang.MetaMethod.doMethodInvoke(MetaMethod.java:328)
58      at org.codehaus.groovy.runtime.metaclass.ClosureMetaClass.invokeMethod(ClosureMetaClass.java:279)
59      at groovy.lang.MetaClassImpl.invokeMethod(MetaClassImpl.java:1009)
60      at groovy.lang.Closure.call(Closure.java:418)
61      at groovy.lang.Closure.call(Closure.java:412)
62      at groovy.lang.Closure.run(Closure.java:500)
63      at java.lang.Thread.run(java.base@18.0.1.1/Thread.java:833)
64
65  Locked ownable synchronizers:
66      - None

```

```

48  at java.lang.invoke.LambdaForm$MH/0x0000000800e27800.guard(java.base@18.0.1.1/LambdaForm$MH
   at java.lang.invoke.DelegatingMethodHandle$Holder.delegate(java.base@18.0.1.1/DelegatingMet
50  at java.lang.invoke.LambdaForm$MH/0x0000000800e27800.guard(java.base@18.0.1.1/LambdaForm$MH
   at java.lang.invoke.Invokers$Holder.linkToCallSite(java.base@18.0.1.1/Invokers$Holder)
52  at ex1$_run_closure2.doCall(ex1.groovy:19)
   at ex1$_run_closure2.doCall(ex1.groovy)
54  at java.lang.invoke.DirectMethodHandle$Holder.invokeSpecial(java.base@18.0.1.1/DirectMethod
   at java.lang.invoke.LambdaForm$MH/0x0000000800c1c800.invoke(java.base@18.0.1.1/LambdaForm$M
56  at java.lang.invoke.Invokers$Holder.invokeExact_MT(java.base@18.0.1.1/Invokers$Holder)
   at jdk.internal.reflect.DirectMethodHandleAccessor.invokeImpl(java.base@18.0.1.1/DirectMeth
58  at jdk.internal.reflect.DirectMethodHandleAccessor.invoke(java.base@18.0.1.1/DirectMethodHa
   at java.lang.reflect.Method.invoke(java.base@18.0.1.1/Method.java:577)
60  at org.codehaus.groovy.reflection.CachedMethod.invoke(CachedMethod.java:343)
   at groovy.lang.MetaMethod.doMethodInvoke(MetaMethod.java:328)
62  at org.codehaus.groovy.runtime.metaclass.ClosureMetaClass.invokeMethod(ClosureMetaClass.jav
   at groovy.lang.MetaClassImpl.invokeMethod(MetaClassImpl.java:1009)
64  at groovy.lang.Closure.call(Closure.java:418)
   at groovy.lang.Closure.call(Closure.java:412)
66  at groovy.lang.Closure.run(Closure.java:500)
   at java.lang.Thread.run(java.base@18.0.1.1/Thread.java:833)
68
   Locked ownable synchronizers:
70   - None

```

One could also make use of online tools that analyse these thread dumps to help identify potential issues. For example, you can try and upload `thread-dump.txt` to this site [fastthread.io](https://fastthread.io) and click on “analyze”.

## 2.4 Classical Synchronization Problems

This section addresses some classical synchronization problems using semaphores.

### 2.4.1 Producers/Consumers

Buffer of size 1, one producer and one consumer. The code below also works if there were multiple producers and multiple consumers.

```

Integer buffer
2
Semaphore consume = new Semaphore(0)
4 Semaphore produce = new Semaphore(1)

6 Thread.start { // Prod
    Random r = new Random()
8     while (true) {
        produce.acquire()
10        buffer = r.nextInt(10000) // produce()
        println "produced "+buffer
12        Thread.sleep(1000)
        consume.release()
14    }
}

16 Thread.start { // Cons
18

```

```

20     while (true) {
21         consume.acquire()
22         println "consumed "+buffer
23         buffer = null // consume(buffer)
24         produce.release()
25     }
26 }

```

Buffer of size  $N$  with one producer and one consumer. Also known as a blocking queue.

```

final int N=10
2 Integer[] buffer = [0] * N

4 Semaphore consume = new Semaphore(0)
5 Semaphore produce = new Semaphore(N)
6 int start = 0
7 int end = 0
8
9 Thread.start { // Prod
10     Random r = new Random()
11     while (true) {
12         produce.acquire()
13         mutexP.acquire()
14         buffer[start] = r.nextInt(10000) // produce()
15         println id+" produced "+buffer[start] + " at index "+start
16         start = (start + 1) % N
17         mutexP.release()
18         consume.release()
19     }
20 }

21 Thread.start { // Cons
22     while (true) {
23         consume.acquire()
24         mutexC.acquire()
25         println id+ " consumed "+buffer[end] + " at index "+end
26         buffer[end] = null // consume(buffer)
27         end = (end + 1) % N
28         mutexC.release()
29         produce.release()
30     }
31 }
32 }

```

Buffer of size  $N$  with multiple producers and multiple consumers.



The static method `currentMethod()` returns a reference to the currently executing thread object. Every thread has a unique id. It may be obtained by using the `getId()` method.

```

final int N=10
2 Integer[] buffer = [0] * N

4 Semaphore consume = new Semaphore(0)
5 Semaphore produce = new Semaphore(N)
6 Semaphore mutexP = new Semaphore(1) // mutex to avoid race conditions on start
7 Semaphore mutexC = new Semaphore(1) // mutex to avoid race conditions on end

```

```

8  int start = 0
   int end = 0
10
12  5.times {
13      Thread.start { // Prod
14          Random r = new Random()
15          while (true) {
16              produce.acquire()
17              mutexP.acquire()
18              buffer[start] = r.nextInt(10000) // produce()
19              println Thread.currentThread().getId()+" produced "+buffer[start] + " at index "+start
20              start = (start + 1) % N
21              mutexP.release()
22              consume.release()
23          }
24      }
25
26  5.times{
27      Thread.start { // Cons
28          while (true) {
29              consume.acquire()
30              mutexC.acquire()
31              println Thread.currentThread().getId()+" consumed "+buffer[end] + " at index "+end
32              buffer[end] = null // consume(buffer)
33              end = (end + 1) % N
34              mutexC.release()
35              produce.release()
36          }
37      }
38  }

```

## 2.4.2 Readers/Writers

## 2.4.3 Barrier Synchronization

```

// One-time use barrier
// Barrier size = N
// Total number of threads in the system = N
4
final int N=3
6  N.times {
7      Thread.start {
8          while (true) {
9              // barrier arrival protocol
10
11              // barrier
12          }
13      }
14  }

```

```

import java.util.concurrent.Semaphore
2  // One-time use barrier
// Barrier size = N
4  // Total number of threads in the system = N

```

```

6  final int N=3
   int t=0
8  Semaphore barrier = new Semaphore(0)
   Semaphore mutex = new Semaphore(1)
10 N.times {
    Thread.start {
12     while (true) {
        // barrier arrival protocol
14     mutex.acquire()
        if (t<N) {
16         t++
            if (t==N) {
18             N.times { barrier.release() }
            }
20     } else {
        barrier.release()
22     }
        mutex.release()
24     // barrier
        barrier.acquire()
26     }
    }
28 }

```

Using cascaded signalling:

```

import java.util.concurrent.Semaphore
2 // One-time use barrier
   // Barrier size = N
4 // Total number of threads in the system = N

6 final int N=3
   int t=0
8 Semaphore barrier = new Semaphore(0)
   Semaphore mutex = new Semaphore(1)
10 N.times {
    Thread.start {
12     while (true) {
        // barrier arrival protocol
14     mutex.acquire()
        if (t<N) {
16         t++
            if (t==N) {
18             barrier.release()
            }
20     }
        mutex.release()
22     // barrier
        barrier.acquire() // Cascaded signalling
24     barrier.release()
    }
26 }
}

```

Cyclic (or reusable) barrier. Failed attempt:

```

1 import java.util.concurrent.Semaphore

```

```

// Cyclic (ie. Reusable) barrier
3 // Barrier size = N
// Total number of threads in the system = N
5
Semaphore mutex = new Semaphore(1)
7 Semaphore barrier = new Semaphore(0)
final int N = 3
9 int t=0
11 N.times {
    Thread.start {
13         while (true) {
            // arrival
15             mutex.acquire()
            t++;
17             if (t==N) {
                N.times { barrier.release()}
19                 t=0 // attempt to reset barrier counter
            }
21             mutex.release()

23             // barrier
            barrier.acquire()
25         }
    }
27 }

```

One easy way to verify that it is incorrect is to count the number of times a thread cycles passed the barrier. Then, notice that some threads can race far ahead of others in terms of the difference in number cycles; this difference can be larger than 1.

A solution follows. We use a second barrier to wait for all threads to fall through the first barrier, thus avoiding any one thread getting ahead of the others.

```

1 import java.util.concurrent.Semaphore
3 // Cyclic (ie. Reusable) barrier
// Barrier size = N
5 // Total number of threads in the system = N
7 Semaphore mutex = new Semaphore(1)
Semaphore barrier = new Semaphore(0)
9 Semaphore barrier2 = new Semaphore(0)
final int N = 3
11 int t=0
13 N.times {
    int id = it
15    Thread.start {
        while (true) {
17            // arrival
            mutex.acquire()
19            t++;
            if (t==N) {
21                N.times { barrier.release() }
            }
23            mutex.release()

```



```

25     // barrier
    barrier.acquire()

27     mutex.acquire()
    t--
29     if (t==0) {
31         N.times { barrier2.release() }
    }
33     mutex.release()

35     barrier2.acquire()
    }
37 }

39 }

1  import java.util.concurrent.Semaphore

3  // Cyclic (ie. Reusable) barrier
  // Barrier size = N
5  // Total number of threads in the system = N

7  Semaphore mutex = new Semaphore(1)
  Semaphore barrier = new Semaphore(0)
9  Semaphore barrier2 = new Semaphore(0)
  final int N = 3
11 int t=0
  int[] c = new int[N]

13 N.times {
15     int id = it
    Thread.start {
17         1000.times { //while (true) {
            // arrival
19             mutex.acquire()
            c[id]++
21             t++
            if (t==N) {
23                 N.times { barrier.release() }
            }
25             mutex.release()

27             // barrier
            println id + " reached barrier. c="+c[id]
29             barrier.acquire()
            println id + " passed barrier. c="+c[id]

31             mutex.acquire()
            t--
33             if (t==0) {
35                 N.times { barrier2.release() }
            }
            mutex.release()

37             barrier2.acquire()
39         }
41     }
}

```

43 }

---

## Chapter 3

# Monitors

A monitor is a program module that encapsulates data and operations and, moreover, guarantees mutual exclusion in the execution of the operations.

Listing 3.1 implements two turnstiles each of which accesses a global counter. The counter is implemented using a monitor. This monitor supports operations `inc()`, `dec()` and `read()`. The `synchronized` qualifier ensures mutual exclusion in the execution of these methods. Every object has a built in lock called an intrinsic lock. When a thread invokes a synchronized method, it automatically acquires the intrinsic lock for that method's object and releases it when the method returns. A thread is said to own the intrinsic lock between the time it has acquired the lock and released the lock. As long as a thread owns an intrinsic lock, no other thread can acquire the same lock. The other thread will block when it attempts to acquire the lock.

```
// Monitor declaration
2  class Counter {
    private int c
4
    public synchronized void inc() {
6        c++
    }
8
    public synchronized void dec() {
10        c--
    }
12
    public synchronized int read() {
14        return c
    }
16 }

18 // Sample use of the monitor
Counter ctr = new Counter()
20
22 P = Thread.start {
    10.times {
24         ctr.inc()
    }
26 }
```

```

Q = Thread.start {
28     10.times {
        ctr.inc()
30     }
    }
32
P.join()
34 Q.join()
println (ctr.read())

```

**Listing 3.1:** Avoiding race conditions on a shared counter using a monitor

### 3.1 A monitor implementing a semaphore

```

1  class Semaphore {
    private int permits
3
    Semaphore(int init) {
5      permits=init
    }
7
    public synchronized void acquire() {
9        while (permits==0) {
            wait()
11       }
        permits--
13    }
15
    public synchronized void release() {
        notify()
17       permits++
    }
19 }
21
Semaphore mutex = new Semaphore(1)
int c=0
23
P = Thread.start {
25     10.times {
        mutex.acquire()
27         c++
        mutex.release()
29     }
    }
31
Q = Thread.start {
33     10.times {
        mutex.acquire()
35         c++
        mutex.release()
37     }
    }
39
P.join()
41 Q.join()
println c

```

This solution is not fair on the threads that are sleeping since an outside thread could “steal” the permit what is made available through a call to `release`. An alternative that is fair in this sense<sup>1</sup> is given in Listing 3.2. A different approach is followed in [Car96].

```

import java.util.concurrent.Semaphore

2
class Semaphore {
4     private int permits;
     private long startWaitingTime=0;
6     private static final long startTime = System.currentTimeMillis();
     private int waiting=0;

8
     Semaphore(int n) {
10         permits=n;
     }

12
     synchronized protected static final long age() {
14         return System.currentTimeMillis() - startTime;
     }

16
     synchronized void acquire() {
18         if (waiting>0 || permits==0) {
             long arrivalTime = age();
20             while (arrivalTime>startWaitingTime || permits==0) {
                 waiting++;
22                 wait();
                 waiting--;
24             }
         }
26         permits--;
     }

28
     synchronized void release() {
30         permits++;
         startWaitingTime = age();
32         notify();
     }
34 }

```

**Listing 3.2:** Fair semaphores

## 3.2 Producers/Consumers

```

class PC {
2     private Object buffer;

4     public synchronized void produce(Object o) {
         while (buffer!=null) {
6             wait()
         }
         buffer = o
8         notifyAll()
     }

```

<sup>1</sup>The idea of using `age` is from [Har98] which uses it in an attempt to propose a fair solution for readers/writers. Unfortunately, the proposed solution is not fair (after an `endWrite` operation, a writer could steal the lock even though there are waiting readers).

```

10     }
12     public synchronized Object consume() {
13         while (buffer==null) {
14             wait()
15         }
16         Object temp = buffer
17         buffer=null
18         notifyAll()
19         return temp
20     }
21 }
22 PC pc = new PC()
23
24 10.times {
25     Thread.start {
26         println (Thread.currentThread().getId()+" consumes")
27         pc.consume()
28     }
29 }
30
31 10.times {
32     Thread.start {
33         println (Thread.currentThread().getId()+" produces")
34         pc.produce((new Random()).nextInt(33))
35     }
36 }

```

Replacing each of the two `notifyAll()` with `notify()` leads to an incorrect solution where one can end up having a producer and consumer both blocked in the wait-set. Hint: C1,C2,P1,P2. This pitfall is called the lost-wakeup problem.

Disadvantages:

- Use multiple condition variables
- Condition variables.

### 3.3 Readers/Writers

Naive solution. Correct but unfair on writers.

```

import java.util.concurrent.locks.*
2
class RW {
4     private int readers;
5     private int writers;
6     static final Lock lock = new ReentrantLock();
7     static final Condition okToRead = lock.newCondition();
8     static final Condition okToWrite = lock.newCondition();
9
10    RW() {
11        readers=0;
12        writers=0;
13    }
14
15    void start_read() {
16        lock.lock();
17        try {

```

```
18     while (writers>0) {
19         okToRead.await();
20     }
21     readers++;
22 } finally {
23     lock.unlock();
24 }
25 }
26
27 void stop_read() {
28     lock.lock();
29     try {
30         readers--;
31         if (readers==0) {
32             okToWrite.signal();
33         }
34     } finally {
35         lock.unlock();
36     }
37 }
38
39 void start_write(Object item) {
40     lock.lock();
41     try {
42         while (readers>0 || writers>0) {
43             okToWrite.await();
44         }
45         writers++;
46     } finally {
47         lock.unlock();
48     }
49 }
50
51 void stop_write() {
52     lock.lock();
53     try {
54         writers--;
55         okToWrite.signal();
56         okToRead.signalAll();
57     } finally {
58         lock.unlock();
59     }
60 }
61
62 }
63
64 RW rw = new RW();
65
66 r = { //R
67     Random r = new Random();
68     rw.start_read();
69     println Thread.currentThread().getId()+" reading..."
70     Thread.sleep(r.nextInt(1000));
71     println Thread.currentThread().getId()+" done reading..."
72
73     rw.stop_read();
74 }
```

```

76 w = { //W
    Random r = new Random();
78    rw.start_write();
    println Thread.currentThread().getId()+" writing..."
80    Thread.sleep(r.nextInt(1000));
    println Thread.currentThread().getId()+" done writing..."
82    rw.stop_write();
    }
84
200.times {
86    Thread.start(r)
    Thread.start(w)
88 }

```

Checking for waiting writers. Places priority on writers but unfair on readers.

```

import java.util.concurrent.locks.*

2
class RW {
4    private int readers;
    private int writers;
6    private int writers_waiting;
    static final Lock lock = new ReentrantLock();
8    static final Condition okToRead = lock.newCondition();
    static final Condition okToWrite = lock.newCondition();

10
    RW() {
12        readers=0;
        writers=0;
14        writers_waiting=0;
    }

16
    void start_read() {
18        lock.lock();
        try {
20            while (writers>0 || writers_waiting>0) {
                okToRead.await();
22            }
            readers++;
24        } finally {
            lock.unlock();
26        }
    }

28
    void stop_read() {
30        lock.lock();
        try {
32            readers--;
            if (readers==0) {
34                okToWrite.signal();
            }
36        } finally {
            lock.unlock();
38        }
    }

40
    void start_write(Object item) {
42        lock.lock();

```



```

44     try {
45         while (readers>0 || writers>0) {
46             writers_waiting++;
47             okToWrite.await();
48             writers_waiting--;
49             writers++;
50         } finally {
51             lock.unlock();
52         }
53     }
54
55     void stop_write() {
56         lock.lock();
57         try {
58             writers--;
59             okToWrite.signal();
60             okToRead.signalAll();
61         } finally {
62             lock.unlock();
63         }
64     }
65 }

```

An attempt at a fair solution to RW is presented in Listing 3.3. One situation that may lead to deadlock is: W1,R1,W2. Another is: R1, W1, R2.

```

1  import java.util.concurrent.locks.*
2
3  class RW {
4      private int readers;
5      private int writers;
6      private int writers_waiting;
7      private int readers_waiting;
8      static final Lock lock = new ReentrantLock();
9      static final Condition okToRead = lock.newCondition();
10     static final Condition okToWrite = lock.newCondition();
11
12     RW() {
13         readers=0;
14         writers=0;
15         writers_waiting=0;
16         readers_waiting=0;
17     }
18
19     void start_read() {
20         lock.lock();
21         try {
22             while (writers>0 || writers_waiting>0) {
23                 readers_waiting++;
24                 okToRead.await();
25                 readers_waiting--;
26             }
27             readers++;
28         } finally {
29             lock.unlock();
30         }
31     }

```

```

33     void stop_read() {
34         lock.lock();
35         try {
36             readers--;
37             if (readers==0) {
38                 okToWrite.signal();
39             }
40         } finally {
41             lock.unlock();
42         }
43     }

44     void start_write(Object item) {
45         lock.lock();
46         try {
47             while (readers>0 || writers>0 || readers_waiting>0) {
48                 writers_waiting++;
49                 okToWrite.await();
50                 writers_waiting--;
51             }
52             writers++;
53         } finally {
54             lock.unlock();
55         }
56     }

57     void stop_write() {
58         lock.lock();
59         try {
60             writers--;
61             okToWrite.signal();
62             okToRead.signalAll();
63         } finally {
64             lock.unlock();
65         }
66     }
67 }
68
69 }

```

**Listing 3.3:** Incorrect attempt at a fair solution to RW; may deadlock

If we replace `stop_write` with the following code, then our solution may deadlock. Hint: Consider W1,R1,W2.

```

1 void stop_write() {
2     lock.lock();
3     try {
4         writers--;
5         if (readers_waiting==0) {
6             okToWrite.signal();
7         } else {
8             okToRead.signalAll();
9         }
10    } finally {
11        lock.unlock();
12    }
13 }

```

## Chapter 4

# Promela

### 4.1 Syntax

We begin with a brief introduction to Promela through a series of examples.

```
1 byte n=0;
3 active proctype P() {
    n=1;
5   printf("P has pid %d. n=%d\n",_pid,n)
   };
7
   active proctype Q() {
9     n=2;
    printf("Q has pid %d. n=%d\n",_pid,n)
11  }
```

Executing Promela code is referred to as a "simulation run of the model".

```
1 $ spin eg.pml
    Q has pid 1. n=2
3   P has pid 0. n=2
2 processes created
```

bash

By default, during simulation runs, SPIN arranges for the output of each active process to appear in a different column: the pid number is used to set the number of tab stops used to indent each new line of output that is produced by a process. You can use the `-t` option to suppress indentation.

```
$ spin -T eg.pml
2 P has pid 0. n=1
  Q has pid 1. n=1
4 2 processes created
```

bash

Semicolon is a separator, not a terminator.

### 4.1.1 Examples involving Loops

```
byte sum=0;
2
active proctype P() {
4   byte i=0;
   do
6     :: i>10 -> break
     :: else ->
8       sum = sum + i;
       i++;
10  od;
12  printf("The sum of the first 10 numbers is %d\n",sum)
}
```

The following example is one of an infinite loop. Run it and note also how SPIN reports overflows errors.

```
1 byte i=0;
3
active proctype P() {
   do
5     :: i++;
     printf("Value of i: %d\n. ",i)
7   od
}
```

An example using a for loop:

```
byte sum=0;
2
active proctype P() {
4   byte i;
   for (i:1..10) {
6     sum = sum + i
   }
8
10  printf("The sum of the first 10 numbers is %d\n",sum)
}
```

### 4.1.2 Expressions as blocking commands

```

byte c=0;
2 finished = 0;
proctype P() {
4   c++;
   finished++
6 }

8 proctype Q() {
   c++;
10  finished++
}

12 init {
14   atomic {
       run P();
16   run Q()
   };
18   finished==2;
   printf("c is %d\n",c)
20 }

```

Equivalently, one may do the following:

```

byte c=0;
2
proctype P() {
4   c++
}

6
proctype Q() {
8   c++
}

10
init {
12   atomic {
       run P();
14   run Q()
   };
16   _nr_pr==1;
   printf("c is %d\n",c)
18 }

```

However, the following variation does not have the expected outcome. When a process terminates, it can only die and make its `_pid` number available for the creation of another process, if and when it has the highest `_pid` number in the system. This means that processes can only die in the reverse order of their creation (in stack order).

```

active proctype P() {
2   printf("A");
}

4
active proctype Q() {
6   printf ("B");
}

8
init {
10  printf("Pr %d",_nr_pr);

```

```

12  _nr_pr==1;
    printf("Done")
}

```

termination.pml

For example, consider what happens if we simulate a run:

```

1  $ spin termination.pml
    A          B          Pr 3          timeout
3  #processes: 3
    3:   proc  2 (:init::1) termination.pml:11 (state 2)
5  3:   proc  1 (Q:1) termination.pml:7 (state 2) <valid end state>
    3:   proc  0 (P:1) termination.pml:3 (state 2) <valid end state>
7  3 processes created

```

bash

It deadlocks at line 11 (`_nr_pr==1`) of the file `termination.pml`. This boolean expression is blocked since processes 0 and 1 cannot terminate until 2 does. If we attempt to verify this program we will obtain an invalid end-state error at line 11.

### 4.1.3 Macros

Semaphores can be modeled in Promela using an `inline` definition. An inline definition works much like a preprocessor macro, in the sense that it just defines a replacement text for a symbolic name, possibly with parameters.

```

byte s=0;

2  inline acquire(s) {
4  atomic {
    s>0 -> s--
6  }
}

8  inline release(s) {
10 s++
}

12 /* AB after CD */
14 proctype P() {
    acquire(s);
16 printf("A");
    printf("B")
18 }

20 proctype Q() {
    printf("C");
22 printf("D");
    release(s)
24 }

26 init {
    atomic {
28 run P();
    run Q()
}

```

```

30 }
    }

```

Problems if you drop the "atomic" in "acquire":

```

1  byte s=1;
   byte c=0;

3

   inline acquire(s) {
5     s>0 -> s--
   }

7

   inline release(s) {
9     s++
   }

11

   proctype P() {
13     acquire(s);
       c++;
15 }

17 proctype Q() {
       acquire(s);
19     c++;
   }

21 /\ AB after CD /\ /\
   /* proctype P() { */
23 /* acquire(s); */
   /* printf("A"); */
25 /* printf("B") */

27 /* } */

29 /* proctype Q() { */

31 /* printf("C"); */
   /* printf("D"); */
33 /* release(s) */
   /* } */

35

   init {
37     atomic {
       run P();
39     run Q()
   }
41     (_nr_pr==1);
       printf("C is %d ",c)
43 }

```

Exercise: would executing lines 7-8 and 18-19 in atomic block avoid deadlock? What about inverting lines 7 and 8 and then placing them in an atomic block (and likewise with lines 18 and 19)?

```
bool wantP = false;
2 bool wantQ = false;
byte cs=0;
4
proctype P() {
6   do
      :: wantP = true;
8     !wantQ;
      cs++;
10    assert (cs==1);
      cs--;
12    wantP=false
   od
14 }

16 proctype Q() {
   do
18     :: wantQ = true;
      !wantP;
20     cs++;
      assert (cs==1);
22     cs--;
      wantQ=false
24   od
   }
26
28 init {
   atomic {
30     run P();
      run Q()
   }
32 }
```

**Figure 4.1:** Attempt III in Promela



```

byte ticket=0;
2 byte mutex=1;

4 inline acquire(sem) {
    atomic {
6       sem>0 -> sem--
    }
8 }

10 inline release(sem) {
    sem++
12 }

14 active [5] proctype Jets() {
    acquire(mutex);
16    acquire(ticket);
    acquire(ticket)
18    release(mutex)
    }

20
22 active [5] proctype Patriot() {
    release(ticket);
    }

```

Figure 4.2: Solution to Bar Problem in Promela

## 4.2 Assertion-Based Model Checking

### 4.2.1 The Bar Problem Revisited

Listing 4.2 presents the solution to the Bar Problem in Promela. We'll verify that this solution is correct in the sense of upholding the problem invariant, namely that there at least two patriots fans for every jets fan. Before doing so, however, let us first run a simulation of this model.

```

1 > spin bar.pml
    timeout
3 #processes: 5
    ticket = 0
5    mutex = 0
    23:   proc 4 (Jets:1) bar.pml:4 (state 4)
7    23:   proc 3 (Jets:1) bar.pml:4 (state 4)
    23:   proc 2 (Jets:1) bar.pml:19 (state 15) <valid end state>
9    23:   proc 1 (Jets:1) bar.pml:19 (state 15) <valid end state>
    23:   proc 0 (Jets:1) bar.pml:4 (state 12)
11 10 processes created

```

bash

The `timeout` indicates that the simulation did not run to completion, it got stuck at a state that is not a valid end state. In other words, it reached a deadlock. From the output above we can see that indeed there are three processes that are deadlocked: 0, 3 and 4. The fact that they are all stuck at line 4 means they are blocked at an `acquire`. Since there are no available permits in `mutex`, clearly processes 3 and 4 are blocked on the `acquire(mutex)` and 0 at the second

acquire(ticket).

A process that terminates must do so after executing its last instruction, otherwise it is said to be in an invalid end state. SPIN checks for this by default. One can insert end state labels to indicate that if execution reaches a certain point and fails to terminate, this should not be considered as an invalid end state. Such valid end state labels must be prefixed with the word `end`. For example, if we replaced the `acquire` operation in 4.2 with the following one:

```

inline acquire(permits) {
2   skip;
end1:
4   atomic {
      permits>0;
6   permits--
      }
8 }

```

then the end states mentioned above are no longer reported as such:

```

> spin bar.pml
2   timeout
#processes: 5
4   ticket = 0
   mutex = 0
6   34:   proc  4 (Jets:1) bar.pml:7 (state 4) <valid end state>
   34:   proc  3 (Jets:1) bar.pml:7 (state 4) <valid end state>
8   34:   proc  2 (Jets:1) bar.pml:22 (state 18) <valid end state>
   34:   proc  1 (Jets:1) bar.pml:7 (state 14) <valid end state>
10  34:   proc  0 (Jets:1) bar.pml:22 (state 18) <valid end state>
10 processes created

```

bash

Let us get back to the task of verifying that the solution is correct. In order to do so we add two counters. Listing 4.2.1 exhibits the updated code.

```

1  byte mutex=1;
   byte ticket=0;
3  byte j=0;
   byte p=0;
5
6  inline acquire(permits) {
7   skip;
end1:
9   atomic {
      permits>0;
11  permits--
      }
13 }
14
15 inline release(permits) {
16   permits++
17 }
18
19 active [5] proctype Jets() {
20
21   acquire(mutex);

```

```

23   acquire(ticket);
    acquire(ticket);
    release(mutex)
25   j++;
    assert (j*2<=p)
27 }

29 active [5] proctype Patriots() {
    release(ticket)
31   p++;
    assert (j*2<=p)
33 }

```

We now verify that our solution is correct.

```

1  $ spin -a bar.pml
   $ gcc -o pan pan.c
3  $ ./pan

5  pan:1: assertion violated ((j*2)<=p) (at depth 34)
   pan: wrote bar.pml.trail

7
   (Spin Version 6.5.1 -- 20 December 2019)
9  Warning: Search not completed
    + Partial Order Reduction

11 Full statespace search for:
13   never claim           - (none specified)
    assertion violations  +
15   acceptance cycles    - (not selected)
    invalid end states   +

17 State-vector 92 byte, depth reached 47, errors: 1
19   18104 states, stored
    18718 states, matched
21   36822 transitions (= stored+matched)
    0 atomic steps
23 hash conflicts:      147 (resolved)

25 Stats on memory usage (in Megabytes):
    2.072   equivalent memory usage for states (stored*(State-vector + overhead))
27   1.071   actual memory usage for states (compression: 51.69%)
    state-vector as stored = 34 byte + 28 byte overhead
29   128.000 memory used for hash table (-w24)
    0.534   memory used for DFS stack (-m10000)
31   129.511 total actual memory usage

33 pan: elapsed time 0.02 seconds
   pan: rate      905200 states/second

```

bash

It seems that this is not the case since an assertion violation is reported. An inspection of the offending trail shows that when the patriots perform a `release(ticket)` but before incrementing

the `p` counter, a jets fan can go in. There are two ways we can fix our code. One is to increment the `p` counter before performing the release. Another one is to perform the release and increment the counter in one atomic block.

### 4.2.2 The MEP Problem

Consider the code for Attempt III from Fig. 4.1. Line 8 and 19 may block. Let us replace these lines with a busy waiting loop, as in the original presentation of Attempt III. The `else` case is important since otherwise the inner `do` loop would be blocked; we want the inner `do` loop to cycle while it waits for the condition to hold.

```

1  bool wantP=false;
2  bool wantQ=false;

4  proctype P() {
5      do
6          :: wantP=true;
7          do
8              :: wantQ==false -> break
9              :: else
10             od;
11         wantP=false
12     od
13 }

14 proctype Q() {
15     do
16         :: wantQ=true;
17         do
18             :: wantP==false -> break
19             :: else
20             od;
21         wantQ=false
22     od
23 }

24 init {
25     atomic {
26         run P();
27         run Q()
28     }
29 }

30

31 bool wantP=false;
32 bool wantQ=false;
33 byte cs=0;

34 proctype P() {
35     do
36         :: wantP=true;
37         do
38             :: wantQ==false -> break
39             :: else
40             od;
41         cs++;
42         assert(cs==1);

```

```

15     cs--;
    wantP=false
16   od
17 }

19 proctype Q() {
    do
21     :: wantQ=true;
        do
23         :: wantP==false -> break
        :: else
25         od;
        cs++;
27         assert(cs==1);
        cs--;
29         wantQ=false
    od
31 }

33 init {
    atomic {
35         run P();
        run Q()
37     }
}

```

```

$ spin -a attemptiii.pml
2 $ gcc -o pan pan.c
$ ./pan

4
(Spin Version 6.5.1 -- 20 December 2019)
6   + Partial Order Reduction

8 Full statespace search for:
    never claim          - (none specified)
10  assertion violations  +
    acceptance  cycles  - (not selected)
12  invalid end states  +

14 State-vector 36 byte, depth reached 14, errors: 0
    21 states, stored
16    21 states, matched
    42 transitions (= stored+matched)
18    1 atomic steps
hash conflicts:          0 (resolved)
20

Stats on memory usage (in Megabytes):
22   0.001   equivalent memory usage for states (stored*(State-vector + overhead))
    0.290   actual memory usage for states
24  128.000   memory used for hash table (-w24)
    0.534   memory used for DFS stack (-m10000)
26  128.730   total actual memory usage

```

```

28 unreachable in proctype P
30     ver.pml:17, state 15, "-end-"
      (1 of 15 states)
32 unreachable in proctype Q
34     ver.pml:31, state 15, "-end-"
      (1 of 15 states)
unreached in init
36     (0 of 4 states)

38 pan: elapsed time 0 seconds

```

bash

### 4.3 Non-Progress Cycles

SPIN can check for some simple liveness properties without the need to use Temporal Logic. An infinite computation that does not include infinitely many occurrences of a progress state is called a non-progress cycle. We illustrate this feature by showing that Attempt III does not enjoy absence of livelock.

Consider

```

byte x=1;

2 active proctype P() {
4     do
6         :: x==1 -> x=2;
          :: x==2 -> x=1;
8     od
}

```

Consider

```

1 byte x=1;

3 active proctype P() {
5     do
7         :: x==1 -> x=2;
          :: x==2 -> progress1: x=1;
9     od
}

```

Consider

```

1 byte x=1;

3 active proctype P() {
5     do
7         :: x==1 -> x=2;
          :: x==2 -> progress1: x=1;
          :: x==2 -> x=1;

```

```

9   od
   }

```

We would like to verify that this attempt at solving the MEP problem does not enjoy absence of livelock. For that we insert progress labels just before entering the CS.

```

bool wantP=false;
2 bool wantQ=false;

4 proctype P() {
    do
        :: wantP=true;
           do
               :: wantQ==false -> break
               :: else
                   od;
           progress1:
               wantP=false
        od
    }

16 proctype Q() {
    do
        :: wantQ=true;
           do
               :: wantP==false -> break
               :: else
                   od;
           progress2:
               wantQ=false
        od
    }

28 init {
    atomic {
30     run P();
    run Q()
    }
32 }

```

Selecting Non-Progress in the drop down list and then verifying, SPIN reports a non-progress cycle:

```

1 2 Q:1 1) wantQ = 1
   Process Statement      wantQ
3 1 P:1 1) wantP = 1      1
   Process Statement      wantP      wantQ
5 2 Q:1 1) else           1          1
   <<<<START OF CYCLE>>>>
7 2 Q:1 1) else           1          1
   1 P:1 1) else           1          1
9 2 Q:1 1) else           1          1
   spin: trail ends after 15 steps

```

spin

“Weak Fairness” should be enabled. Weak fairness means that each statement that becomes enabled and remains enabled thereafter will eventually be scheduled. Consider the example below [?]:

```

1 byte x=0;
2
3 active proctype P() {
4     do
5         :: true -> x = 1 - x;
6     od
7 }
8
9 active proctype Q() {
10    do
11        :: true -> progress1: x = 1 - x;
12    od
13 }

```

It is possible that Q makes no progress if Q is never scheduled for execution. Weak fairness guarantees that it eventually will. Verify this in SPIN by first enabling weak fairness and then disabling it. In the former case no errors are reported, but in the latter a non-progress cycle is reported:

```

1 0 P:1 1) 1
  <<<<<START OF CYCLE>>>>>
3 0 P:1 1) x = (1-x)
  Process Statement      x
5 0 P:1 1) 1             1
  0 P:1 1) x = (1-x)     1
7 0 P:1 1) 1             0

```

spin

Consider the code for Attempt IV

```

1 bool wantP = false, wantQ = false;
2
3 active proctype P() {
4     do
5         :: wantP = true;
6         do
7             :: wantQ -> wantP = false; wantP = true
8             :: else -> break
9         od;
10        wantP = false
11    od
12 }
13
14 active proctype Q() {
15    do
16        :: wantQ = true;
17        do
18            :: wantP -> wantQ = false; wantQ = true
19            :: else -> break
20        od;
21    wantQ = false
22 }

```



```

23 }
    od

```

We know that it does not enjoy freedom from starvation. Freedom from starvation would mean that both P and Q enter their CS infinitely often. We can verify that it does not enjoy freedom from starvation by inserting a progress label in the critical section of P, selecting Non-Progress in the drop down list and then verifying.

```

1  bool wantP = false, wantQ = false;

3  active proctype P() {
    do
5     :: wantP = true;
        do
7         :: wantQ -> wantP = false; wantP = true
            :: else -> break
        od;
    progress1:
11     wantP = false
        od
13 }

15 active proctype Q() {
    do
17     :: wantQ = true;
        do
19         :: wantP -> wantQ = false; wantQ = true
            :: else -> break
        od;
    progress2:
23     wantQ = false
        od
25 }

```

Here is the output from SPIN

```

1  1 Q:1  1)  wantQ = 1
    Process Statement      wantQ
3  1 Q:1  1)  else
    1 Q:1  1)  wantQ = 0    1
5  0 P:1  1)  wantP = 1    0
    Process Statement      wantP      wantQ
7  1 Q:1  1)  wantQ = 1    1          0
    1 Q:1  1)  wantP
    0 P:1  1)  wantQ        1          1
    <<<<START OF CYCLE>>>>
11 1 Q:1  1)  wantQ = 0    1          1
    1 Q:1  1)  wantQ = 1    1          0
13 1 Q:1  1)  wantP
    0 P:1  1)  wantP = 0    1          1
15 1 Q:1  1)  wantQ = 0    0          1
    1 Q:1  1)  wantQ = 1    0          0
17 1 Q:1  1)  else
    0 P:1  1)  wantP = 1    0          1
19 0 P:1  1)  wantQ
    1 Q:1  1)  wantQ = 0    1          1
21 0 P:1  1)  wantP = 0    1          0
    1 Q:1  1)  wantQ = 1    0          0

```

```

23 1 Q:1 1) else 0 1
    1 Q:1 1) wantQ = 0 0 1
25 0 P:1 1) wantP = 1 0 0
    1 Q:1 1) wantQ = 1 1 0
27 1 Q:1 1) wantP 1 1
    Process Statement wantP wantQ
29 0 P:1 1) wantQ 1 1
    spin: trail ends after 50 steps

```

### 4.3.1 The Zoo Problem Revisited

Consider the Zoo Problem discussed in Exercise ??:

In a zoo there is a common feeding area for exotic mice and exotic felines. The feeding area has a common feeding lot that holds up to 2 animals. The feeding area can be used by both mice and felines, but cannot be used by mice and felines at the same time for obvious reasons.

A solution in Promela is given in Listing 4.3.

**Exercise 4.3.1.** Show that if lines are removed, then deadlock is possible. Explain the deadlock situation that can arise.

**Exercise 4.3.2.** Show, using assertions, that there cannot be felines feeding, if there are mice feeding and, likewise, there cannot be mice feeding, if there are felines feeding.

**Exercise 4.3.3.** You are asked to show that there can be at most two mice and at most two felines in the feeding lot. Since you already know that both cannot be feeding at the same time, you propose the following assertions (only the portion of code that is modified is shown) in the Mouse process:

```

    acquire(feedinglot);
2  // access feeding lot
    assert (mice<3);
4  release(feedinglot);

```

Similarly for the Feline process:

```

    acquire(feedinglot);
2  // use feeding lot
    assert (felines<3);
4  release(feedinglot);

```

Unfortunately, when you verify this with Spin, it reports an assertion violation.

1. Explain why.
2. Replace the assertions with new ones that address the problem correctly.

```

byte mice = 0;
2 byte felines = 0;
byte mutexMice = 1;
4 byte mutexFelines = 1;
byte feedinglot = 2;
6 byte mutex = 1;

8 inline acquire(sem) {
    atomic {
10     sem>0;
        sem--
12     }
}

14 inline release(sem) {
16     sem++
}

18 active [3] proctype Mouse() {
20     acquire(mutex);
    acquire(mutexMice);
22     mice++;
    if
24     :: mice==1 -> acquire(mutexFelines);
        :: else -> skip;
26     fi
    release(mutexMice);
28     release(mutex);

30     acquire(feedinglot);
    // access feeding lot
32     release(feedinglot);

34     acquire(mutexMice);
    mice--;
36     if
        :: mice==0 -> release(mutexFelines);
        :: else -> skip;
38     fi
    release(mutexMice);
40 }

42 active [3] proctype Feline() {
44     acquire(mutex);
    acquire(mutexFelines);
46     felines++;
    if
48     :: felines==1 -> acquire(mutexMice);
        :: else -> skip;
50     fi
    release(mutexFelines);
52     release(mutex);

54     acquire(feedinglot);
    // use feeding lot
56     release(feedinglot);

58     acquire(mutexFelines);
    felines--;
60     if
        :: felines==0 -> release(mutexMice);
        :: else -> skip;
62     fi
    release(mutexFelines);
64 }
66 }

```

CP Notes

v0.015

Figure 4.3: Zoo Problem in Promela



## Chapter 5

# Solution to Selected Exercises

### Section ??

Answer 5.0.1 (Exercise ??). *jj*

### Section ??



# Bibliography

- [Car96] Tom Cargill. Specific notification for java thread synchronization. [www.dre.vanderbilt.edu/%7Eeschmidt/PDF/specific-notification.pdf](http://www.dre.vanderbilt.edu/%7Eeschmidt/PDF/specific-notification.pdf), 1996.
- [Har98] Stephen Hartley. Concurrent Programming: The Java Programming Language. Oxford University Press, 1998.